

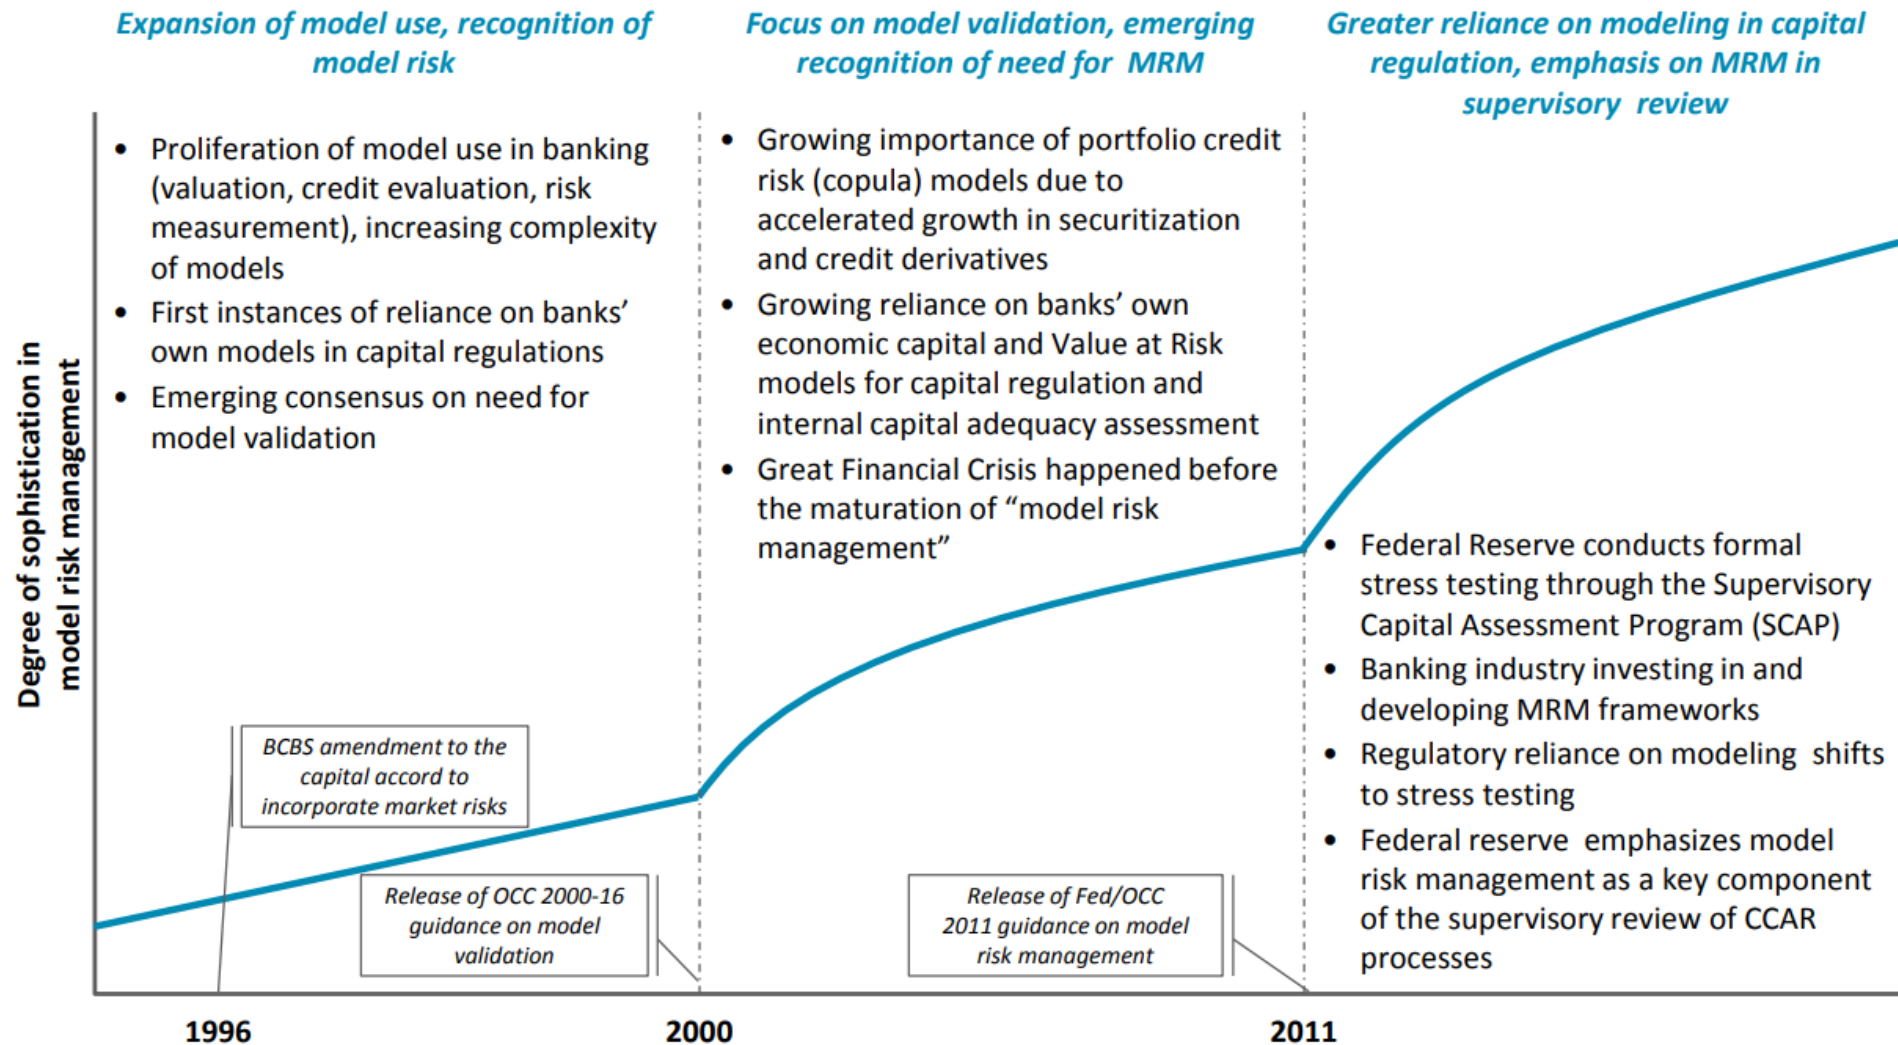
RESPONSIBLE AI

Week 5: AI Regulations

TODAY'S OBJECTIVES


- What precedents are there for AI regulation? Why does AI need special regulation?
- What does the impossibility theorem mean for designing laws around AI? Where is humanity today in terms of AI regulation? What will the regulatory landscape look like as you graduate into data science roles in the workforce?
- How can regulations require AI to be fair and unbiased when those definitions aren't standardized or compatible?

HISTORY OF REGULATIONS AROUND QUANTITATIVE MODELS



WHY DO WE NEED AI-SPECIFIC REGULATION?

- Scale
- Intersection and application across specialized fields
 - Cybersecurity and privacy
 - Medicine
 - Human capital / employment law
 - Content moderation
 - Housing
 - Ethics
- Reliance on data
 - Bias within different stages of the AI model lifecycle
- Complexity
- Opacity

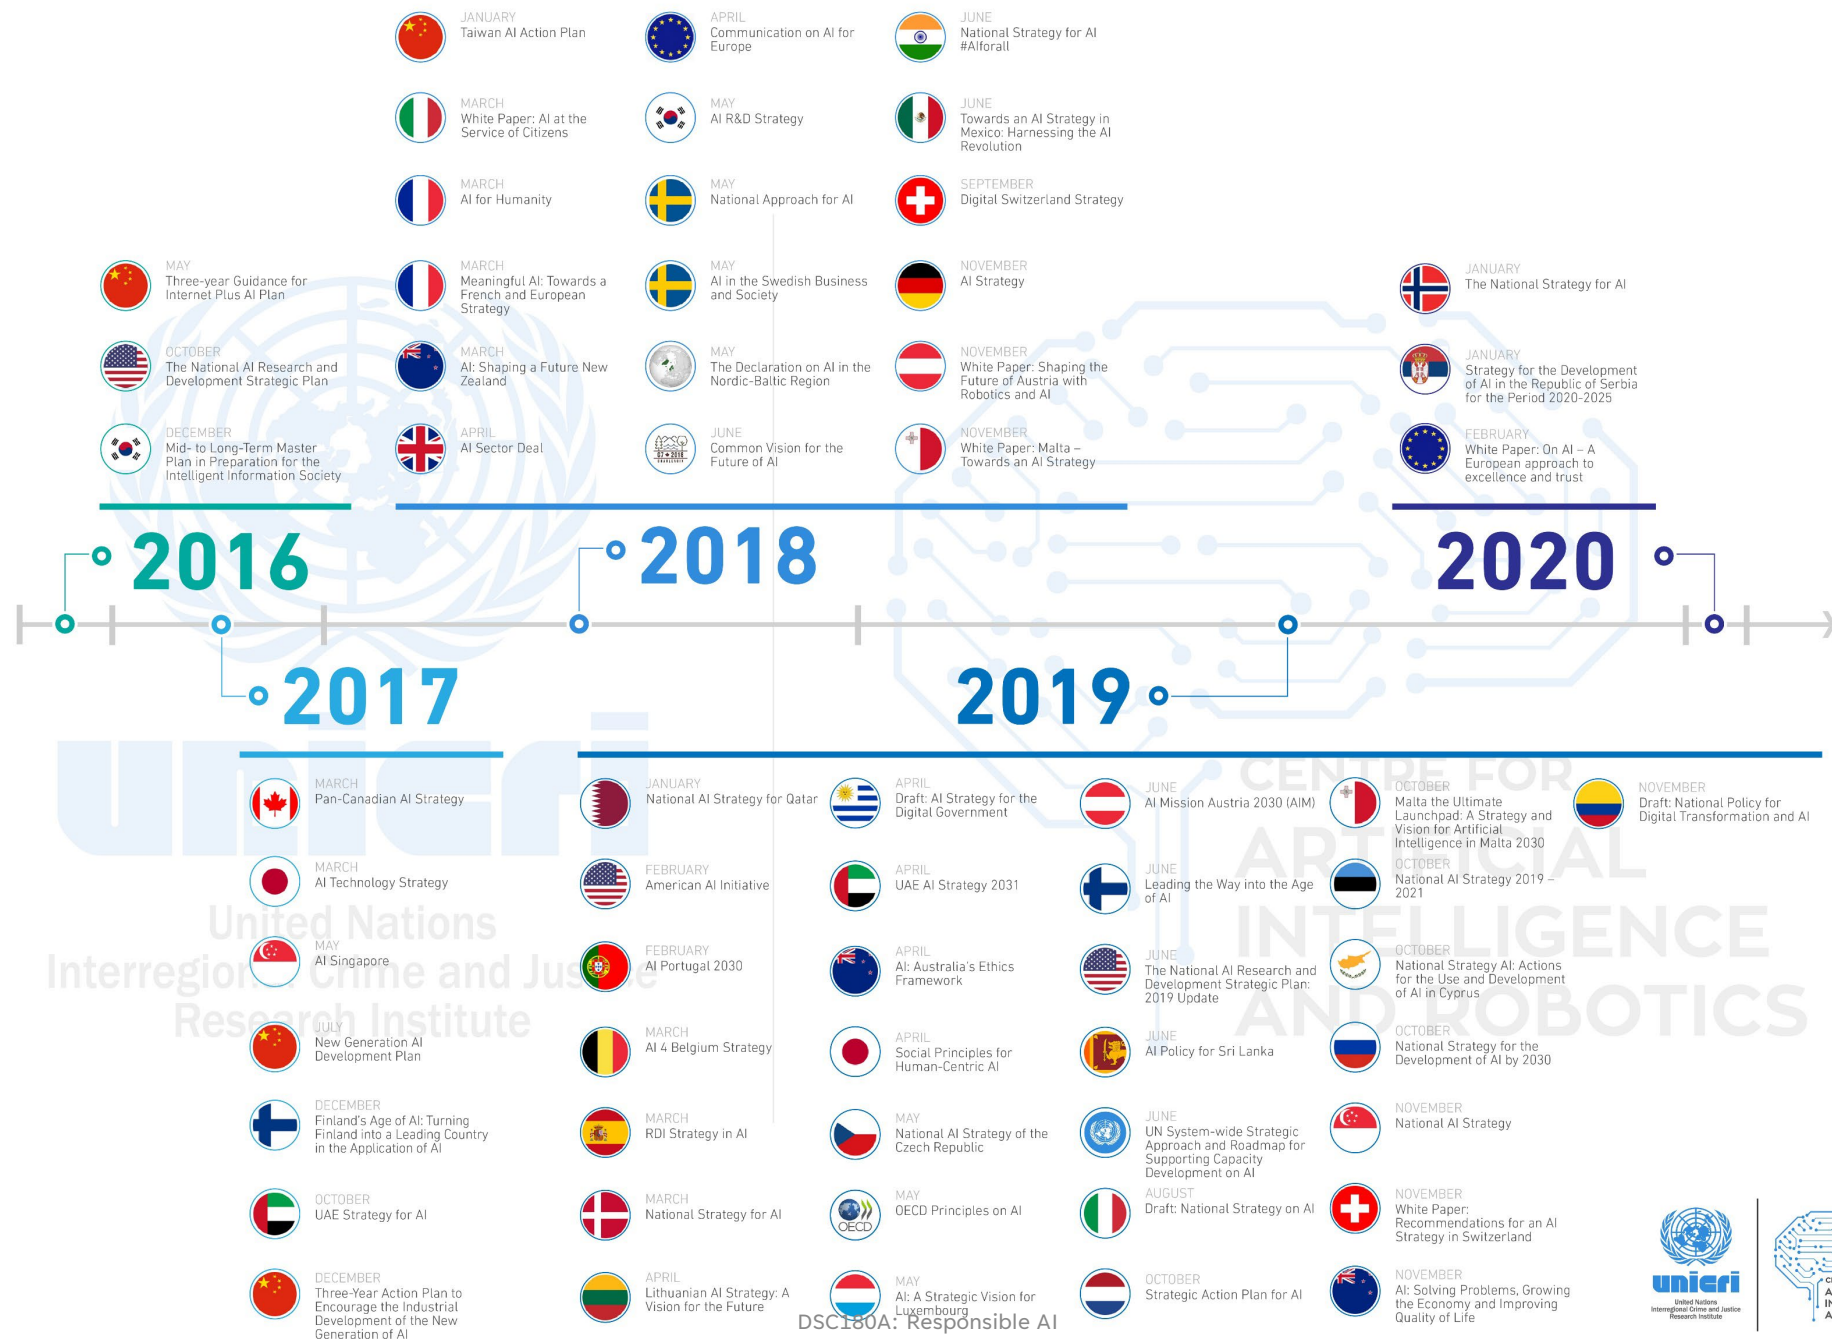


READING PRESENTATION #1

*Big Data's Disparate Impact
(Barocas and Selbst)*

Challenges & Limitations to AI Regulations

- **Impossibility Theorem:** How to choose and balance definitions of fairness?
- **(Lack of) Standardization:** There is no “source of truth” for AI Regulations – how will these regulations be consistently assessed and implemented?
- **False Sense of Security:** Will users “let their guard down” if they believe their AI is regulated?
- **Legal Issues:** Bias Mitigation as Affirmative Action?
 - Equal-protection doctrine may threaten algorithmic fairness. Equal protection doctrine prevents government entities from treating similarly situated individuals differently (“anticlassification”). **This approach prohibits making decisions based on protected attributes** (i.e. hiring someone because of their race or gender). However, leading bias mitigation approaches often **purposefully use protected attributes to promote fair or unbiased outcomes in their algorithms**
 - There is a high level of legal scrutiny required to bypass anticlassification, which is why affirmative action cases are so frequently brought to the courts
 - The courts are likely to consider leading bias mitigation techniques a type of affirmative action, which subjects them to strict scrutiny (= harder to get passed) of the law and even potentially renders them illegal



AI Policy Efforts in the US

FEBRUARY 11, 2019

Executive Order 13859: Maintaining American Leadership in Artificial Intelligence

Announced the American AI Initiative
*"The United States must foster **public trust and confidence in AI technologies** and protect civil liberties, privacy, and American values in their application in order to fully realize the potential of AI technologies for the American people."*

FEBRUARY 26, 2020

American AI Initiative: Year One Annual Report

Provided a summary of progress and long-term vision for the American AI Initiative, emphasizing the need to **"embrace trustworthy AI for government services and missions."**

NOVEMBER 17, 2020

Office of Management and Budget (OMB) Memorandum M-21-06: Guidance for Regulation of Artificial Intelligence Applications

Provided guidance to Federal agencies, including considerations for reducing barriers to AI development and adoption
*"The government's regulatory and non-regulatory approaches to AI should contribute to public trust in AI by **promoting reliable, robust, and trustworthy AI applications.**"*

DECEMBER 3, 2020

Executive Order 13960: Promoting the Use of Trustworthy Artificial Intelligence in the Federal Government

Outlined a set of principles and actions to accelerate trustworthy AI use and development

JANUARY 1, 2021

National AI Initiative Act of 2020

*"The mission of the National AI Initiative is to ensure continued U.S. leadership in AI research and development, **lead the world in the development and use of trustworthy AI** in the public and private sectors..."*

OCTOBER 4, 2022

Blueprint for the AI Bill of Rights

Outlines five protections Americans should have in the AI age:

1. Safe and Effective Systems
2. Algorithmic Discrimination Protection
3. Data Privacy
4. Notice and Explanation, and
5. Human Alternatives, Consideration, and Fallback

What would binding AI regulation look like?

BRIEFING

EU Legislation in Progress



Artificial intelligence act

OVERVIEW

The European Commission unveiled a new proposal for an EU regulatory framework on artificial intelligence (AI) in April 2021. The draft AI act is the first ever attempt to enact a horizontal regulation of AI. The proposed legal framework focuses on the specific utilisation of AI systems and associated risks. The Commission proposes to establish a technology-neutral definition of AI systems in EU law and to lay down a classification for AI systems with different requirements and obligations tailored on a 'risk-based approach'. Some AI systems presenting 'unacceptable' risks would be prohibited. A wide range of 'high-risk' AI systems would be authorised, but subject to a set of requirements and obligations to gain access to the EU market. Those AI systems presenting only 'limited risk' would be subject to very light transparency obligations. While generally supporting the Commission's proposal, stakeholders and experts call for a number of amendments, including revising the definition of AI systems, broadening the list of prohibited AI systems, strengthening enforcement and redress mechanisms and ensuring proper democratic oversight of the design and implementation of EU AI regulation.

Trustworthy AI | Regulatory Landscape for Human Capital

Federal labor laws and protections have been found in courts to extend to AI-based decisions...

... to protect candidates regardless of age, gender, race, religion, nationality, disabilities, etc. ...

... throughout the sourcing, contacting, screening, and interview of potential candidates...

... all the way down to the state and local level.

Title VII of the Civil Rights Act of 1964 (Title VII)

- a federal law that protects employees and applicants against discrimination based on certain specified characteristics such as race, color, national origin, sex, and religion
- prohibits discrimination based on disparate treatment and/or disparate impact
- a court could find that an employer faces the same liability for a program exhibiting the unconscious bias of its programmer as it would if the programmer had made the hiring decision him or herself, based on that bias.

Age Discrimination in Employment Act (ADEA)

- forbids age discrimination against people who are age 40 or older. It does not protect workers under the age of 40, although some states have laws that protect younger workers from age discrimination

Americans with Disabilities Act (ADA)

- An A.I.-hiring practice could also implicate the Americans with Disabilities Act (“ADA”) if an algorithm discerns an applicant’s physical disability, mental health, or clinical diagnosis, all of which are forbidden inquiries in pre-employment candidate assessments. The ADA Amendments Act of 2008 broadened the statutory definition of “disability,” increasing the scope of individuals whom the ADA protects.

Equal Employment Opportunity Commission (EEOC)

- enforces disability discrimination laws with respect to employers in the private sector and the federal government
- has issued guidance qualifying the expanded list of personality disorders identified in the psychiatric literature as protected mental impairments
- has already investigated at least two instances of alleged A.I. bias, and has made clear that employers using A.I. hiring practices could face liability for any unintended discrimination.
- in 2017, found reasonable cause to believe an employer violated the ADEA by advertising on Facebook for a position within its company and “limiting the audience for their advertisement to younger applicants.”

Uniformed Services Employment and Reemployment Rights Act (USERRA)

- Maintains reemployment, antidiscrimination, and antiretaliation rights for members of uniformed services

State and Local Legislation

- New York City
 - Local law 144 of 2021 prohibits employers from using an automated employment decision tool unless such tool has been subject to a bias audit within one year of the use of the tool. (Laws goes live in Jan 2023).
- Illinois (Artificial Intelligence Video Interview Act)
 - imposes strict limitations on employers who use A.I. to analyze candidate video interviews. Employers must a) notify applicants that A.I. will be used in their interviews; b) obtain consent to use A.I.; c) explain to applicants how the A.I. works and what characteristics the A.I. will track in relation to the position; d) limit sharing of the video interview; and e) comply with requests to destroy the video within 30 days
- Maryland
 - the prohibition of facial/voice recognition software in interviews

DISCUSSION

The NYC Department of Consumer and Worker Protection recently released additional rules and clarification around the NYC hiring law.

New York City Department of Consumer and Worker Protection

Notice of Public Hearing and Opportunity to Comment on Proposed Rules

What are we proposing? The Department of Consumer and Worker Protection (“DCWP” or “Department”) is proposing to add rules to implement new legislation related to the use of automated employment decision tools. The proposed rules would clarify the requirements for the use of automated employment decision tools within New York City, the notices to employees and candidates for employment regarding the use of the tool, the bias audit for the tool, and the required published results of the bias audit.

Follow the link provided in [chat](#) and discuss these **three** questions in your breakout rooms:

- **What metric of fairness is the Department proposing?** (ctrl+f ‘impact ratio’, ‘selection rate’, and ‘bias audit’)
- **What are the strengths and weaknesses of this approach?** Does this go far enough?
- Unlike familiar financial audits, which have been refined over decades of accounting experience, the AI audit process is new and without clearly established guidelines:
 - ... the Surveillance Technology Oversight Project, New York Civil Liberties Union and other organizations has noted the lack of standards for bias audits
 - ... businesses and their service providers are grappling with how to comply with New York City’s mandate for AI audits used in hiring

What does challenges could this pose for different stakeholders (think through the different stakeholders we discussed in previous lectures!)?

11

Make sure to nominate someone to represent your group and share your thoughts with the room when we reconvene!

FOR NEXT WEEK

- Complete next week's **readings**
 - If you signed up to present **Ethical Machine Learning in Health Care (Chen et al. 2021)** come prepared to present next week and submit your presentation to Gradescope by 10 AM PT, Friday, November 4th
- Submit your answers to next week's participation questions on Gradescope by 10 PM PT, Thursday, November 3rd
- **Replication Part #2:** Notebook and writeup on model development and fairness metric assessments
 - Primary contact for replication project: Nandita Rahman (nanrahman@deloitte.com)
 - Office hours: Mondays 1-2pm PST
- First replication project checkpoint is this weekend: submit EDA notebook